

Behavioral Analysis of the Agent-Based Community Grid Solution for the Large Hadron Collider beauty Experiment

Daniela Remenska^{1,3}, Tim A.C. Willemse², Henri Bal¹, Kees Verstoep¹, Wan Fokkink¹, and Jeff Templon³

¹Dept. of Computer Science, VU University Amsterdam, The Netherlands

²Dept. of Computer Science, TU Eindhoven, The Netherlands

³NIKHEF, Amsterdam, The Netherlands

Abstract—DIRAC (Distributed Infrastructure with Remote Agent Control) is the Grid solution designed to support production activities as well as user data analysis for the Large Hadron Collider beauty experiment. It consists of cooperating distributed services and a plethora of light-weight agents delivering the workload to the Grid resources. Services accept requests from agents and running jobs, while agents actively fulfill specific goals. Services maintain database back-ends to store dynamic state information of entities such as jobs, queues, or staging requests. Agents continuously check for changes in the service states, and react to these accordingly. The logic of each agent is rather simple; the main source of complexity lies in their cooperation. These agents run concurrently, and communicate using the services' databases as a shared memory for synchronizing the state transitions. Despite the effort invested in making DIRAC reliable, entities occasionally get into inconsistent states. Tracing and fixing such behaviors is difficult, given the inherent parallelism among the distributed components, and the size of the implementation.

In this paper we present an analysis of DIRAC with mCRL2, process algebra with data. We have reverse engineered two critical and related DIRAC subsystems, and subsequently modeled their behavior with the mCRL2 toolset. This enabled us to easily locate race conditions and livelocks which were confirmed to occur in the real system. We further formalized and verified several behavioral properties of the two modeled subsystems.

Keywords: DIRAC, service-oriented architecture, agents, stager, mCRL2, model checking, process algebra

I. INTRODUCTION

The Large Hadron Collider beauty (LHCb) experiment [1] is one of the four large experiments conducted on the Large Hadron Collider (LHC) accelerator, built by the European Organization for Nuclear Research (CERN). Immense amounts of data are produced at the LHC accelerator, and subsequently processed by physics groups and individuals worldwide. The sheer size of the experiment is the motivation behind the adoption of the Grid computing paradigm. The Grid storage and computing resources for the LHCb experiment are distributed across several institutes in Europe. To cope with the complexity of processing the vast amount of data, a complete Grid solution, called DIRAC (Distributed Infrastructure with Remote Agent Control) [2], [3], has been designed and developed for the LHCb community.

DIRAC forms a layer between the LHCb community of users and the heterogeneous Grid resources, in order to allow for optimal and reliable usage of these resources. It consists of many cooperating distributed services and light-weight agents which deliver workload to the resources. The logic of each individual component is relatively simple; the overall system complexity emerges from the cooperation among them. Namely, these agents run concurrently, and communicate using the services' databases as a shared memory (blackboard paradigm [4]) for synchronizing state transitions of various entities.

Although much effort is invested in making DIRAC reliable, entities occasionally get into inconsistent states, leading to a potential loss of efficiency in both resource usage and manpower. Debugging and fixing the root of such encountered behaviors becomes a formidable mission due to multiple factors: the inherent parallelism present among the system components deployed on different physical machines, the size of the implementation (around 150000 lines of Python code), and the distributed knowledge of different subsystems within the collaboration.

In this paper we propose the use of more rigorous (formal) methods for improving software quality. Model checking [5] is one such technique for analysis of an abstract model of a system, and verification of certain system properties of interest. Unlike conventional testing, it allows full control over the execution of parallel processes and also supports automated exhaustive state-space exploration.

We used the mCRL2 language [6] and toolset [7] to model the behavior of two critical and related DIRAC components: the Workload Management (WMS) and the Storage Management System (SMS). Based on Algebra of Communicating Processes (ACP) [8], mCRL2 is able to deal with generic data types as well as user-defined functions for data transformation. This makes it particularly suitable for modeling the data manipulations made by DIRAC's agents. Visualizing the state space and replaying scenarios with the toolkit's simulator enabled us to gain insight into the system behavior, incrementally improve the model, and to already detect critical race-conditions and livelocks, which were confirmed to occur in

the real system. Some of them were a result of simple coding bugs; others unveiled more elementary design problems. We further formulated, formalized and verified several general and application-specific properties.

The idea of modeling existing systems using formal techniques is as such not new. Earlier studies ([9],[10],[11],[12],[13],[14],[15]) mostly focused on modeling and verifying hardware or communication protocols, since the formal languages and tools at hand were not sufficiently mature to cope with more complex data-intensive distributed systems. More recently, success stories on modeling real-life concurrent systems with data have been reported ([16],[17],[18],[19],[20]). In [18] the authors have implemented a tool for automatic translation of the SystemC language into mCRL2 statements. This greatly simplifies the analysis, but has so far been feasible only when the language of implementation is domain-specific, or alternatively, a reasonably small subset of a general-purpose language is considered for translation. The only exception in this respect is the Java Pathfinder tool [19] used to find deadlocks and other behavioral properties in Java software systems developed by NASA.

We believe that the challenges and results of this work are unique in a number of aspects. First, to the best of our knowledge, the code-base and the number of concurrent components engaged in providing DIRACs functionality considerably outnumber previous industrial cases. Second, the choice of Python as implementation platform has lead to prevailing usage of dynamic structures (whose types and sizes are determined at runtime) throughout DIRAC, challenging the transition to an abstract formal representation. We have nevertheless established general guidelines on extracting a model outline from the implementation. Third, analysis of this kind is typically performed after a problem has already surfaced in the real system, as a means to understand the events which lead to it and test for possible solutions. We managed to stumble on an actual bug at the same time it was observed in practice, which increased our confidence in the soundness of the model.

The paper is organized as follows. Section 2 introduces the architecture of DIRAC, focusing on the two studied subsystems. Section 3 gives a brief overview of the mCRL2 language, and describes our approach to abstracting and modeling the behavior of these subsystems. Section 4 presents the analysis with the mCRL2 toolset and the issues detected. Section 5 concludes and discusses future work.

II. DIRAC: A COMMUNITY GRID SOLUTION

A. Architecture Overview

The development of DIRAC started in 2002 as a system for production of simulation data that would serve to verify theory, aspects of the LHCb detector design, as well as to optimize algorithms. It gradually evolved into a complete community Grid solution for data and job management, based on a general-purpose framework that can be reused by other communities besides LHCb. Today, it covers all major LHCb tasks starting with the raw data transfer from the experiments

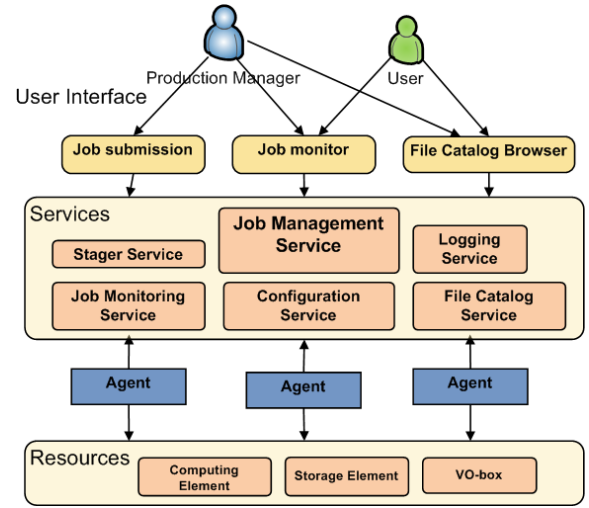


Fig. 1. DIRAC Architecture overview

detector to the Grid storage, several steps of data processing, up to the final user analysis. Python was chosen as the implementation language, since it enables rapid prototyping and development of new features. DIRAC follows the Service Oriented Architecture (SOA) paradigm, accompanied by a network of lightweight distributed agents which animate the system. Its main components are depicted in Fig. 1.

The *services* are passive components that react to requests from their clients, possibly soliciting other services in order to fulfill the requests. They run as permanent processes deployed on a number of high-availability hosts (VO-boxes) at CERN, and persist the dynamic system state information in database repositories. The user interfaces, agents or running jobs can act as clients placing the requests to DIRACs services.

Agents are active components that fulfill a limited number of specific system functions. They can run in different environments, depending on their mission. Some are deployed close to the corresponding services, while others run on the Grid worker nodes. Examples of the later are the so-called Pilot Agents, part of the Workload Management System explained in the following section. All DIRAC agents repeat the same logic in each iteration cycle: they observe for changes in the service states, and react accordingly by initiating actions (like job submission or data transfer) which may update the states of various system entities.

Resources are software abstractions of the underlying heterogeneous Grid computing and storage resources allocated to LHCb, providing a uniform interface for access. The physical resources are controlled by the site managers and made available through middleware services such as gLite [21].

The DIRAC functionality is exposed to users and developers through a rich set of command-line tools forming the DIRAC API, complemented by a Web portal for visual monitoring the system behavior and controlling the ongoing tasks. Both the Web and command-line *interfaces* ensure secure system access using X509 certificates.




Fig. 2. DIRAC Workload Management System [22]

B. Workload Management System

The driving force of DIRAC is the Workload Management System (WMS). Taking into account that the distributed computing environment is intrinsically unstable, a *Pilot Job paradigm*, illustrated in Fig. 2, was chosen as an efficient way to implement a pull scheduling mechanism. Pilot jobs are simply resource reservation processes without an actual payload defined apriori. They are submitted to the Grid worker nodes with the aim of checking the sanity of the operational environment just before pulling and executing the real payload. This hides the fragility of the underlying resources and increases the job success rate, from the perspective of end users. Centralized *Task Queues* contain all the pending jobs, organized in groups according to their requirements. This enables efficient application of job priorities, enforcing fair-share policies and quotas, which are decided within the LHCb community, without the need of any effort from the Grid sites.

The basic flowchart describing the evolution of a jobs states is depicted in Fig. 3. After submission through the *Job Manager* (Received), the complete job description is placed in the DIRAC job repository (the Job DB). Before jobs become eligible for execution, a chain of optimizer agents check and prioritize them in queues, utilizing the parameters information from the Job DB. The *Job Sanity Agent* can meaningfully fail jobs with impossible requirements, and prevent unnecessary submission to the Grid resources. The

Input Data Agent performs a resolution of possible target computing resources based on the input data requirements. If the requested data resides on tape storage, the *Job Scheduling Agent* will pass the control to a specialized *Stager* service (part of the SMS explained in the next section), before placing the job a *Task Queue* (Waiting). Based on the complete list of pending payloads, a specialized *Task Queue Director* submits pilots to the computing resources via the gLite WMS. After a *Matcher* service pulls the most suitable payload for a pilot (Matched), a *Job Wrapper* object is created on the worker node, responsible for retrieving the input sandbox, performing software availability checks, executing the actual payload on the worker node (Running), and finally uploading any output data necessary (Done or Failed). The wrapper can catch any failure exit state reported by the running physics applications. At the same time, a *Watchdog* process is instantiated to monitor the behavior of the Job Wrapper and send heartbeat signals to the monitoring service. It can also take actions in case resources are soon to be exhausted, the payload stalls, or a management command for killing the payload is received.

Although the Grid storage resources are limited, it is essential to keep all data collected throughout the experiments run. Tape backends provide a reliable and cheap solution for data storage. The additional workflow step necessary for input data files residing on tape is carried out inside the Storage Management System (SMS).

C. Storage Management System

The DIRAC SMS provides the logic for pre-staging files from tape to a disk cache frontend, before a job is able to process them. Smooth functioning of this system is essential for production activities which involve reprocessing of older data with improved physics software, and happens typically several times per year.

A simplified view of the system is shown in Fig. 4. The workflow is initiated with the Job Scheduling agent detecting that a job is assigned to process files only available on tape storage. It sends a request for staging (i.e., creating a cached replica) to the *Storage Manager Handler* service along with the list of files and a callback method to be invoked when the request has been processed. The Storage Management DB is immediately populated with records which are processed by a sequence of agents in an organized fashion. The relevant tables in the SMS DB are the *Tasks* and *CacheReplicas*, whose entities maintain a state observed and updated by these agents. Tasks maintain general information about every job requesting a service from the SMS. The details about every file (i.e., the Storage Element (SE) where it resides, the size, checksum, number of tasks that requested it), are kept in the *CacheReplicas* table. Other auxiliary tables maintain the relationship between these entities. The processing begins with the *Request Preparation Agent*. It selects all the New entries for files, checks whether they are registered in a special logical file catalog, and retrieves their metadata. In case of problematic catalog entries, it can update the status of the *CacheReplicas* and the related *Tasks* entries to Failed. For non-problematic

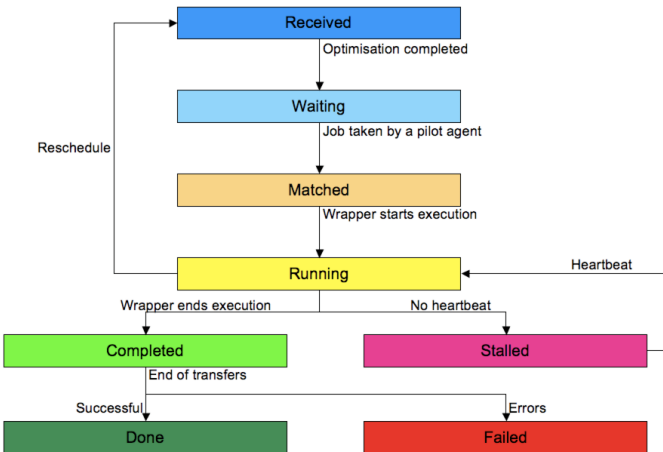


Fig. 3. Job state machine [23]

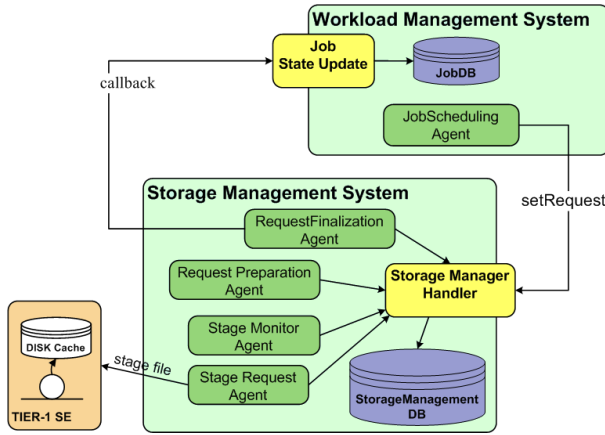


Fig. 4. DIRAC Storage Management System

files it carries out transition to Waiting. The *StageRequest Agent* is responsible for placing the actual staging requests for all Waiting entries, via dedicated storage middleware that communicates with the tape backends. These requests are grouped by SE prior to submission, and carry information about the requested (pin) lifetime of the replicas to be cached. If certain pathologies are discovered (i.e., lost, unavailable or zero-sized files on tape), it can update the corresponding entries to Failed in a similar manner. Otherwise, they are promoted to StageSubmitted. The agent responsible for monitoring the status of submitted requests is the *StageMonitor Agent*. It achieves this by interrogating the storage middleware to see if the StageSubmitted files are successfully replicated on disk cache. In case of success, the CacheReplicas and their corresponding Tasks entries are updated to Staged. Various circumstances of tape or middleware misbehavior can also fail the staging requests. The last one in the chain is the *RequestFinalization Agent*. The tasks which are in their final states (Staged or Failed) are cleared from the database, and callbacks are performed to the WMS, which effectively wakes up the corresponding jobs. If there are no more associated tasks for particular replicas, the respective CacheReplicas entries are also removed.

This subsystem is still in evolutionary phase. In multiple instances, tasks or replicas have become stuck, effectively blocking the progress of jobs. Tracing back the sequence of events which led to the inconsistent states is non-trivial. To temporarily alleviate such problems, the status of these entries is typically reset to the initial New state manually, so agents can re-process them from scratch. Occasionally, error messages are reported from unsuccessful attempts of the SMS service to update the state of a non-existent table entry.

REFERENCES

[1] Large Hadron Collider beauty experiment. [Online]. Available: <http://lhcb-public.web.cern.ch/lhcb-public>

[2] A. Tsaregorodtsev, M. Bargiotti, N. Brook, A. R. Casajus, G. Castellani, P. Charpentier, C. Cioffi, J. Closier, R. G. Diaz, G. Kuznetsov, Y. Y. Li, R. Nandakumar, R. Paterson, A. Smith, M. S. Miguelez, and S. G. Jimenez, "DIRAC: A community grid solution," in *Proc. Computing in High Energy Physics (CHEP 2007)*, vol. 119, no. 2. IOP Publishing.

[3] A. Smith and A. Tsaregorodtsev, "DIRAC: Reliable data management for LHCb," in *Proc. Computing in High Energy Physics (CHEP 2007)*, vol. 119, no. 2. IOP Publishing.

[4] J. McManus and W. Bynum, "Design and analysis techniques for concurrent blackboard systems," *Systems, Man and Cybernetics*, vol. 26, pp. 669–680, Nov. 1996.

[5] J.F.Groote and T.A.C.Willemse, "Model-checking processes with data," *Science of Computer Programming*, vol. 56, no. 3, pp. 251–273, 2005.

[6] J.F.Groote, A. Mathijssen, M. Reniers, Y. Usenko, and M. V. Weerdenburg, "The formal specification language mCRL2," 2007.

[7] J.F.Groote *et al.*, "The mCRL2 toolset," Eindhoven University of Technology, Tech. Rep., 2008.

[8] T. Basten, J. Baeten, and M.A.Reniers, *Process Algebra: Equational Theories of Communicating Processes*. NY, USA: Cambridge University Press, 2009.

[9] Y. He and R. Janicki, "Verifying protocols by model checking: a case study of the wireless application protocol and the model checker SPIN," in *Proc. CASCON 2004*. IBM Press.

[10] K. Palmisano, "Verification of the session management protocol," Royal Institute of Technology, Masters Thesis, 2006.

[11] G. Holzmann, *Design and validation of computer protocols*. Upper Saddle River, NJ, USA: Prentice-Hall Inc., 1990.

[12] W. Fokkink, *Modelling Distributed Systems: Protocol Verification with μ CRL*, 2nd ed. Springer, 2011.

[13] T. Ball and S. Rajamani, "The SLAM toolkit," in *Proc. Computer Aided Verification (CAV 2001)*. Springer-Verlag, pp. 260–264.

[14] B. Badban, W. Fokkink, J. Groote, J. Pang, and J. V. Pol, "Verification of a sliding window protocol in μ CRL and PVS," *Formal Aspects of Computing*, vol. 17, no. 3, October 2005.

[15] S. Islam, M. Sqalli, and S. Khan, "Modeling and formal verification of DHCP using SPIN," *IJCSA*, May 2006.

[16] J. Hwon, V. Kuster, and T. Willemse, "Analysing the control software of the Compact Muon Solenoid experiment at the Large Hadron Collider," in press.

[17] D. Bošnački, A. Mathijssen, and Y. Usenko, "Behavioural analysis of an I2C linux driver," *Proc. FMICS 2009*.

[18] H. Hojjat, M. Mousavi, and M. Sirjani, "Formal analysis of SystemC designs in process algebra," *Fundamenta Informaticae*, pp. 19–42, 2011.

[19] V. Visser and P. Mehlitz, "Model checking programs with Java PathFinder," in *Proc. SPIN (2005)*.

[20] B. Ploeger, "Analysis of ACS using mCRL2," Technische Universiteit Eindhoven, CS-Report 09-11, 2009.

[21] glite - Lightweight Middleware for Grid Computing. [Online]. Available: <http://glite.cern.ch>

[22] A. Casajus, R. G. Diaz, S. Paterson, and A. Tsaregorodtsev, "DIRAC pilot framework and the DIRAC Workload Management System," in *Proc. Computing in High Energy Physics (CHEP 2010)*. IOP Publishing.

[23] Grid production shifter guide for LHCb. [Online]. Available: <https://twiki.cern.ch/twiki/bin/view/LHCb/ProductionShifterGuide>